

# Milestone 1: Index Construction Report

## Algorithms and Data Structures Developer Option

This report presents the analytics for the inverted index constructed using the **Algorithms and Data Structures Developer** approach. This option is required for CS and SE students.

### Approach Specifications:

- **Corpus:** All ICS web pages (developer.zip) - ~56,000 pages
- **Index Storage:** File system (no databases) - multiple partial indexes merged
- **Memory:** Cannot hold entire index in memory - uses disk-based offloading
- **Offloading:** Index offloaded to disk at least 3 times during construction
- **Search Response Time:** Target  $\leq 300\text{ms}$  (ideally  $\leq 100\text{ms}$ )
- **Programming Level:** Advanced - efficient data structures and file access

### Index Specifications:

- Tokens: All alphanumeric sequences
- Stop words: Not used (all words indexed)
- Stemming: Porter stemming algorithm
- Important words: Words in bold, headings (h1-h3), and titles are marked
- Term frequency: Calculated for each token in each document

## Index Analytics

Metric	Value
Number of Indexed Documents	1,212
Number of Unique Tokens	13,126
Total Size of Index on Disk	14438.72 KB
Partial Indexes Created (Offloads)	4

### ■ Requirement Verification:

The indexer successfully created 4 partial indexes during construction, meeting the requirement of at least 3 offloads to disk.

## Implementation Details

### Disk-Based Indexing Architecture:

The indexer uses a memory-efficient approach that periodically offloads the in-memory index to disk as

partial index files. This ensures the index can be built for datasets of any size without running out of memory.

**Processing Workflow:**

1. Documents are processed and added to an in-memory index chunk
2. When memory limit is reached, the current index is saved as a partial index file
3. In-memory index is cleared and processing continues
4. Steps 1-3 repeat until all documents are processed
5. All partial indexes are loaded and merged into a single final index
6. Final merged index is saved to disk
7. Partial index files are cleaned up

**File Organization:**

- **inverted\_index.json:** Final merged inverted index
- **doc\_mapping.json:** URL to document ID mappings
- **partial\_indexes/:** Temporary directory containing partial indexes during construction

**Memory Management:**

Memory usage is bounded by the chunk size (typically 300-5,000 documents per chunk). This ensures the indexer can handle very large datasets (56,000+ pages) without exceeding available memory. The chunk size is automatically calculated to ensure at least 3 offloads occur during construction.

**Search Component Requirements (Milestone 2):**

The search component must read postings from disk without loading the entire index into memory. This requires efficient disk-based lookups and data structures optimized for fast query response times.